

# COMP336 — Big Data

## Week 7 Lecture 1: Finding Similar Items

Diego Mollá

Department of Computer Science  
Macquarie University

COMP348 2018H1

# Who am I



Diego Mollá Aliod

Email: [diego.molla-aliod@mq.edu.au](mailto:diego.molla-aliod@mq.edu.au)

Office: E6A331

Webpage:

<http://web.science.mq.edu.au/~diego/>

Research interests: Text processing,  
machine learning: question answering,  
named entity recognition, text  
summarisation.

# What We've Seen So Far

## Amin Beheshti

- ① Data and Big Data
- ② Organizing Big Data
- ③ Curating Big Data
- ④ Processing Big Data: Cloud computing
- ⑤ Processing Big Data: MapReduce (I)
- ⑥ Processing Big Data: MapReduce (II)

These are techniques for organising, curating, and **general techniques** for processing Big Data.

# What We're Yet to See

## Diego Molla

- 7 Finding Similar Items
- 8 Mining Data Streams
- 9 Link Analysis
- 10 Frequent Itemsets
- 11 Large-scale Machine Learning (I)
- 12 Large-scale Machine Learning (II)

These are **specific techniques** for particular problems that require the processing of big data.

# Programme

- 1 Mining of Massive Datasets
- 2 Shingling and Minhashing
- 3 Locality-Sensitive Hashing

## Reading

- Leskovec, Rajaraman, Ullman (2014): Mining of Massive Datasets, Chapter 3. <http://www.mmids.org/>

# Programme

- 1 Mining of Massive Datasets
  - Data Mining
  - Finding Similar Items
- 2 Shingling and Minhashing
- 3 Locality-Sensitive Hashing

# Programme

- 1 Mining of Massive Datasets
  - Data Mining
  - Finding Similar Items
- 2 Shingling and Minhashing
- 3 Locality-Sensitive Hashing

# Mining of Big Data

- MapReduce is a technique that manages parallel processing of big data.
- It can be very effective for tasks that require massive parallelisation.
- But there is a trade-off between computation and resources.
  - We can reduce computation time by increasing the number of computation nodes.
  - If the number of computation nodes is limited, MapReduce may take long to complete.
- Sometimes we need specific methods for particular tasks.



# Programme

- 1 Mining of Massive Datasets
  - Data Mining
  - Finding Similar Items
- 2 Shingling and Minhashing
- 3 Locality-Sensitive Hashing


# Applications of Near-Neighbour Search

- Find plagiarised documents.
- Detect mirror pages.
- Find news articles from the same source.
- Collaborative Filtering.
  - On-line purchases.
  - Movie ratings.

# Jaccard Similarity

- The Jaccard similarity of two sets  $A$  and  $B$  is:

$$\frac{|A \cap B|}{|A \cup B|}$$



- We want to find efficient methods to:
  - 1 Represent documents as sets.
  - 2 Compute the Jaccard similarity between  $A$  and  $B$ , where  $A$  and  $B$  can be large sets.
  - 3 Find the pairs with highest Jaccard similarity in a large collection of sets.

# Programme

- 1 Mining of Massive Datasets
- 2 **Shingling and Minhashing**
  - Shingling
  - Minhashing
- 3 Locality-Sensitive Hashing

# Programme

- 1 Mining of Massive Datasets
- 2 Shingling and Minhashing
  - Shingling
  - Minhashing
- 3 Locality-Sensitive Hashing

# Shingling

- A document can be characterised by the set of words appearing in it.
- But this would ignore ordering of words.
- Shingles are an alternative that can represent documents ...
  - ... as sets of something ...
  - ... that account for word ordering (of some sort).

# $k$ -Shingles

## $k$ -shingle

Define a  $k$ -shingle for a document to be any substring of length  $k$  found within the document.

## Example

A document  $D$  is the string  $abcdabd$ . If  $k = 2$ , then the set of 2-shingles for  $D$  is:

$$\{ab, bc, cd, da, bd\}$$

Note that the substring  $ab$  appears twice within  $D$  but appears only once as a shingle.

# Matrix Representation of Sets

A (sparse) matrix can represent a collection of sets by using the following convention:

- Each row indicates a possible element.
- Each column represents a set.
- A value of 1 in row  $r$  and column  $c$  indicates that the element  $r$  is in set  $c$ .

## Example

Element	$S_1$	$S_2$	$S_3$	$S_4$
$a$	1	0	0	1
$b$	0	0	1	0
$c$	0	1	0	1
$d$	1	0	1	1
$e$	0	0	1	0

This matrix represents the following sets:

- 1  $S_1 = \{a, d\}$
- 2  $S_2 = \{c\}$
- 3  $S_3 = \{b, d, e\}$
- 4  $S_4 = \{a, c, d\}$



# How to Represent Shingles Efficiently?


- Assuming that there are 32 characters, then there are  $32^9$  different 9-shingles.
- A matrix with  $32^9$  rows will not fit in memory of a regular computer.
- But most shingles will never occur, so we can represent a shingle using a **hash function**.

## Hashing Shingles

- Set a hash function that maps a shingle to a number within a limited range, e.g. 0 to  $2^{32} - 1$ .
- The shingle is represented by its hash number.
- We can now represent all shingles in 4 bytes.

# What is a Hash Function?

- A hash function maps an item (string, number, list, etc) to a number within a predefined range.
- A good hash function will minimise the number of **collisions**.
  - Two different items that map to the same value.
- Hashes are used in programming languages to speed up the search in a list.



# Exercise I

Why not using, say, 4-shingles?

- 4-shingles may not have the power to differentiate documents.
- Represent two documents using 4-shingles and compute the Jaccard difference.
- Represent the same two documents using 9-shingles and compute the Jaccard difference.

## Exercise II

### Why using 9-shingles hashed down to 4 bytes?

- The set of 9-shingles hashed down to 4 bytes takes the same space as the set of 4-shingles.
- But we use most of the hash space to represent 9-shingles and we underuse much of the space to represent 4-shingles.
- Represent the same two documents using 9-shingles hashed down to 4 bytes and compute the Jaccard difference.

### Trade-off

Different shingles may map to the same hash number (hash collisions) so we trade accuracy for scalability.

# Programme

- 1 Mining of Massive Datasets
- 2 Shingling and Minhashing
  - Shingling
  - Minhashing
- 3 Locality-Sensitive Hashing

# Motivation for Minhashing and Locality Sensitive Hashing

- Even if we can represent a document by its  $k$ -shingles, how do we find all similar documents?
- Imagine you have  $N = 1$  million documents.
- Naively, if we compute pairwise Jaccard similarities for every pair of docs, you would need  $N \times (N - 1)/2 \approx 5 \times 10^{11}$  comparisons.
- At  $10^5$  secs per day and  $10^6$  comparisons per sec, it would take 5 days.
- For  $N = 10$  million it would take more than a year!

# Minhashing

- To minhash a set represented by a column of the characteristic matrix, pick a permutation of the rows.
- The minhash value of a column is the first row, in the permuted order, in which the column has a 1.

## Example

Element	$S_1$	$S_2$	$S_3$	$S_4$
$b$	0	0	1	0
$e$	0	0	1	0
$a$	1	0	0	1
$d$	1	0	1	1
$c$	0	1	0	1

$$h(S_1) = a, h(S_2) = c, h(S_3) = b, h(S_4) = a$$

# Minhashing and Jaccard Similarity

## Minhashing and Jaccard Similarity

The probability that the minhash function for a random permutation of rows produces the same value for two sets equals the Jaccard similarity of those sets.

- Minhashing is a good hash function for computing the Jaccard similarity.
- If the Jaccard similarity between two sets  $C_1$  and  $C_2$  is high, then with high probability  $h(C_1) = h(C_2)$ .
- If the Jaccard similarity between two sets  $C_1$  and  $C_2$  is low, then with high probability  $h(C_1) \neq h(C_2)$ .



# Minhash Signatures

- 1 Use several (e.g. 100) independent permutations of the characteristic matrix.
- 2 Compute the minhashing according to each permutation.
- 3 Now we have 100 independent signatures of each document.
- 4 These 100 signatures form the rows of a matrix called *signature matrix*.
- 5 This new signature matrix has many less rows than the original characteristic matrix!

# Computing Minhash Signatures

- It is not practical to permute a large characteristic matrix explicitly.
- Instead, we will simulate a random permutation by defining a random **hash function** that maps row numbers to their new row positions.

## Procedure

- 1 Initialise  $SIG(i, c) \leftarrow \infty$ ,  $i = 1 \cdots n$ , for all  $c$
- 2 Scan row  $r = 1 \cdots M$  of the characteristic matrix.
- 3 Compute  $h_1(r), h_2(r), \dots, h_n(r)$
- 4 For each column  $c$  do the following:
  - 1 If  $c$  has 0 in row  $r$ , do nothing.
  - 2 If  $c$  has 1 in row  $r$ ,  
 $SIG(i, c) \leftarrow \min(SIG(i, c), h_i(r))$ ,  $i = 1 \cdots n$

# Example

Let's define two hash functions (on the right of the characteristic matrix below).

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1) \% 5$	$(3x + 1) \% 5$
$a$	1	0	0	1	1	1
$b$	0	0	1	0	2	4
$c$	0	1	0	1	3	2
$d$	1	0	1	1	4	0
$e$	0	0	1	0	0	3

# Example - Initialisation

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1)\%5$	$(3x + 1)\%5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	$\infty$	$\infty$	$\infty$	$\infty$
$h_2$	$\infty$	$\infty$	$\infty$	$\infty$

## Example - After Scanning Row 0

$\Rightarrow$

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1)\%5$	$(3x + 1)\%5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	1	$\infty$	$\infty$	1
$h_2$	1	$\infty$	$\infty$	1

## Example - After Scanning Row 1

$\Rightarrow$

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1)\%5$	$(3x + 1)\%5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	1	$\infty$	2	1
$h_2$	1	$\infty$	4	1

## Example - After Scanning Row 2

$\Rightarrow$

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1) \% 5$	$(3x + 1) \% 5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	1	3	2	1
$h_2$	1	2	4	1

## Example - After Scanning Row 3

$\Rightarrow$

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1)\%5$	$(3x + 1)\%5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
4	0	0	1	0	0	3

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	1	3	2	1
$h_2$	0	2	0	0



## Example - After Scanning Row 4

Element	$S_1$	$S_2$	$S_3$	$S_4$	$(x + 1)\%5$	$(3x + 1)\%5$
0	1	0	0	1	1	1
1	0	0	1	0	2	4
2	0	1	0	1	3	2
3	1	0	1	1	4	0
⇒ 4	0	0	1	0	0	3

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	1	3	0	1
$h_2$	0	2	0	0

## Exercise: Compare Jaccard Similarities

Fill the following tables showing the pairwise Jaccard similarities using the characteristic matrix and the signature matrix.

Characteristic Matrix

	$S_1$	$S_2$	$S_3$	$S_4$
$S_1$		0	$1/4$	$2/3$
$S_2$				
$S_3$				
$S_4$				

Signature Matrix

	$S_1$	$S_2$	$S_3$	$S_4$
$S_1$		0	$1/2$	1
$S_2$				
$S_3$				
$S_4$				

# Programme

- 1 Mining of Massive Datasets
- 2 Shingling and Minhashing
- 3 Locality-Sensitive Hashing**

## What we Know So Far

- How to represent documents as sets.
- How to compress the sets to small signatures.
- How to use the small signatures to approximate the Jaccard Similarity.

But

How do we efficiently find the most similar documents?

# Locality-Sensitive Hashing

- It is very time-consuming to find the expected similarity of all pairs of documents.
    - Would you use MapReduce for this...?
  - But often we only need to find the pairs of documents that are most similar.
  - We will study an approach that can be used for documents represented by shingle-sets and then minhashed to short signatures.
- ⇒ Read Section 3.6 of the textbook if you want to learn the general theory of locality-sensitive hashing.

## LSH for Minhashing — The Key Idea

- Hash items several times, in such a way that similar items are more likely to be hashed to the same bucket than dissimilar items are.
- We then consider any pairs that hashed to the same bucket for any of the hashings to be candidate pairs.
- The key is to find hash functions that reduce errors:
  - False positives: Dissimilar pairs that hash to the same bucket.
  - False negatives: Similar pairs that do not hash to the same bucket.
- What is worse, a false positive or a false negative...?

# LSH Using the Signature Matrix

- We want to find hash functions that put similar columns of the signature matrix in the same buckets.
- Remember that we are using the Jaccard similarity metric, so columns that share many values in their rows are more similar.

## Big Idea

Define hash functions for **bands** of the signature matrix.

# Approach

- 1 Divide signature matrix into  $b$  bands of  $r$  rows.
- 2 For each band, hash its portion of each column to a hash table with  $k$  buckets.
  - Make  $k$  as large as possible.
  - We want to avoid collisions which would lead to false positives.
  - We are using hashing as a quick method to compare column fragments.
- 3 Candidate column pairs are those that hash to the same bucket for at least one band.
- 4 Tune  $b$  and  $r$  to catch most similar pairs, but few non-similar pairs.



## Example — Signature Matrix

This is an artificial example! The signature matrix would have many more rows and columns.

	$S_1$	$S_2$	$S_3$	$S_4$
$r_1$	1	0	1	1
$r_2$	2	1	0	2
$r_3$	1	2	2	0
$r_4$	0	1	1	2
$r_5$	3	0	3	1
$r_6$	1	0	1	0

## Example — After Dividing Matrix

This is an artificial example! The signature matrix would have many more rows and columns.

	$S_1$	$S_2$	$S_3$	$S_4$
$r_1$	1	0	1	1
$r_2$	2	1	0	2
<hr/>				
$r_3$	1	2	2	0
$r_4$	0	1	1	2
<hr/>				
$r_5$	3	0	3	1
$r_6$	1	0	1	0

## Example — After Hashing Columns

This is an artificial example! The signature matrix would have many more rows and columns.

	$S_1$	$S_2$	$S_3$	$S_4$
$r_1$	1	0	1	1
$r_2$	2	1	0	2
$r_3$	1	2	2	0
$r_4$	0	1	1	2
$r_5$	3	0	3	1
$r_6$	1	0	1	0

## Example — Final Candidate Pairs

This is an artificial example! The signature matrix would have many more rows and columns.

	$S_1$	$S_2$	$S_3$	$S_4$
$r_1$	1	0	1	1
$r_2$	2	1	0	2
$r_3$	1	2	2	0
$r_4$	0	1	1	2
$r_5$	3	0	3	1
$r_6$	1	0	1	0


Candidate pairs are:  $(S_1, S_4)$ ,  $(S_2, S_3)$ ,  $(S_1, S_3)$

## Analysis of This Technique

Suppose that we use  $b$  bands of  $r$  rows each, and suppose that a particular pair of documents has Jaccard similarity  $s$ .

- The probability that the signatures agree in all rows of one particular band is  $s^r$ .
- The probability that the signatures do not agree in at least one row of a particular band is  $1 - s^r$ .
- The probability that the signatures do not agree in all rows of any of the bands is  $(1 - s^r)^b$ .
- The probability that the signatures agree in all rows of at least one band is  $1 - (1 - s^r)^b$ .
- The formula  $(1/b)^{1/r}$  approximates the value of  $s$  where the probability is 0.5.

## Picking $r$ and $b$ : the S-curve



# Locality Search Hashing: The Complete Procedure

- 1 Pick a value of  $k$  and construct the set of  $k$ -shingles of each document.
  - Optionally, hash the  $k$ -shingles to shorter bucket numbers.
- 2 Sort the document-shingle pairs to order them by shingle.
- 3 Pick a length  $n$  for the minhash signature matrix. Compute the minhash signature matrix.
- 4 Choose a threshold  $t$  that decides the jaccard similarity of two documents to be to be considered “similar”. Select values of  $b$  and  $r$  so that  $t$  is approximately  $(1/b)^{1/r}$ .
- 5 Construct candidate pairs by applying LSH.
- 6 Examine each candidate pairs more in detail.

Which of these steps can be parallelised?

## Take-home Messages

- **Shingling**: A way to convert documents to sets.
- **Minhashing**: A way to convert large sets to short signatures.
- **Locality-Sensitive Hashing (LSH)**: An efficient way to detect pairs of similar documents from a large collection.
- Both minhashing and LSH trade efficiency for accuracy.



# What's Next

## Week 8

- RECESS: 16-29 April.
- Assignment 2 deadline: Workshop week 9.
- Topic week 8: Mining Data Streams.