

COMP336 — Big Data

Week 8 Lecture 1: Mining Data Streams

Diego Mollá

Department of Computer Science
Macquarie University

COMP348 2018H1

Programme

- 1 Data Streams
- 2 Sampling and Filtering
- 3 Counting Elements in a Stream

Reading

- Leskovec, Rajaraman, Ullman (2014): Mining of Massive Datasets, Chapter 4. <http://www.mmds.org/>

Programme

- 1 Data Streams
- 2 Sampling and Filtering
- 3 Counting Elements in a Stream

Characteristics of Data Streams

- Velocity:** Data may arrive faster than we can process it.
- Volume:** Accumulated data might not fit in the available storage space. We can think of data as **infinite**.
- Variety:** Data may change in time. Data that happened some time ago might not be relevant any more. We can think of data as **non-stationary**.
- (This is not the standard meaning of variety...)
- We still need to handle the “classic” issue of variety: we may need to handle multiple streams at once.

Examples of Data Streams

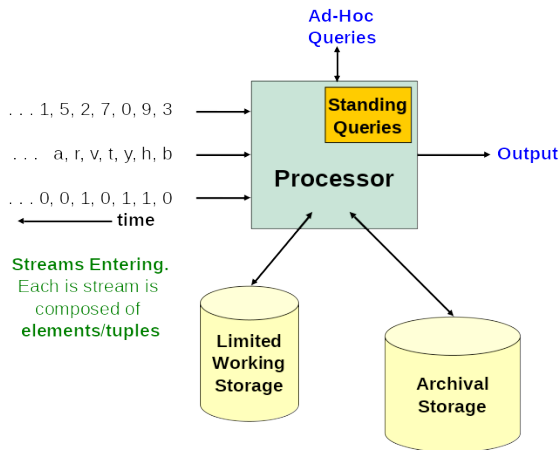
Image Data: Surveillance cameras, satellite imagery, ...

Sensor data: Temperature, GPS coordinates, heart rate, ...

Internet and Web Traffic :

- Search queries;
- Posts from Twitter, Facebook, ...
- IP packets;
- Clicks.

The Stream Model



Streams Entering.
Each stream is
composed of
elements/tuples

<http://www.mmds.org/>

Storage in the Stream Model

Archival Storage

- Large storage for archival purposes.
- We assume it is not possible to answer queries from the archival store.
- Can be used only under special circumstances using time-consuming retrieval processes.

Working Store

- Holds summaries or parts of streams.
- Can be used for answering queries.
- Might be in disk or in main memory.
- Cannot store all the data from all the streams.

Storage in the Stream Model

Archival Storage

- Large storage for archival purposes.
- We assume it is not possible to answer queries from the archival store.
- Can be used only under special circumstances using time-consuming retrieval processes.

Working Store

- Holds summaries or parts of streams.
- Can be used for answering queries.
- Might be in disk or in main memory.
- Cannot store all the data from all the streams.

Types of Queries

Standing Queries

- Queries that are always performed on the data.
- In a sense, these are queries that are permanently executing.
- Since these queries are known in advance, it is fairly easy to design efficient storage and query processes to handle them.

Ad-Hoc Queries

- Queries that are not known in advance.
- These queries are created, for example, by a user or operator.
- We need to find a way to query the current state of the stream.

Types of Queries

Standing Queries


- Queries that are always performed on the data.
- In a sense, these are queries that are permanently executing.
- Since these queries are known in advance, it is fairly easy to design efficient storage and query processes to handle them.

Ad-Hoc Queries

- Queries that are not known in advance.
- These queries are created, for example, by a user or operator.
- We need to find a way to query the current state of the stream.

Examples of Standing Queries

Example: Ocean Surface Temperature Sensor



- ❶ Alert when the temperature exceeds 25 degrees centigrade.
- ❷ Average the 24 most recent readings.
- ❸ Maximum temperature ever recorded.
- ❹ Average temperature.

Question

What information do we need to keep in the working storage to answer each of these standing queries?

Examples of Working Storage Needs

Q1: Alert when the temperature exceeds 25°C

- No information required (unless we want to allow a threshold input by the user)

Q2: Average the 24 most recent readings

- 24 variables, one per reading.

Q3: Maximum temperature ever recorded

- 1 variable with the value of the maximum so far.

Q4: Average temperature

- 1 variable with the value of the sum of readings so far.
- 1 variable that counts the number of readings so far.

Examples of Working Storage Needs

Q1: Alert when the temperature exceeds 25°C

- No information required (unless we want to allow a threshold input by the user)

Q2: Average the 24 most recent readings

- 24 variables, one per reading.

Q3: Maximum temperature ever recorded

- 1 variable with the value of the maximum so far.

Q4: Average temperature

- 1 variable with the value of the sum of readings so far.
- 1 variable that counts the number of readings so far.

Examples of Working Storage Needs

Q1: Alert when the temperature exceeds 25°C

- No information required (unless we want to allow a threshold input by the user)

Q2: Average the 24 most recent readings

- 24 variables, one per reading.

Q3: Maximum temperature ever recorded

- 1 variable with the value of the maximum so far.

Q4: Average temperature

- 1 variable with the value of the sum of readings so far.
- 1 variable that counts the number of readings so far.

Examples of Working Storage Needs

Q1: Alert when the temperature exceeds 25°C

- No information required (unless we want to allow a threshold input by the user)

Q2: Average the 24 most recent readings

- 24 variables, one per reading.

Q3: Maximum temperature ever recorded

- 1 variable with the value of the maximum so far.

Q4: Average temperature

- 1 variable with the value of the sum of readings so far.
- 1 variable that counts the number of readings so far.

Question: An effective way to compute the average temperature

Q4: Average temperature

If we keep the sum of readings so far we may have problems with data overflow (the sum may exceed the capacity of storage)

- 1 How serious is this problem?
- 2 How could we fix this problem?

Examples of Ad-hoc Queries

Example: Web Site

- 1 What were the unique users in the past month?
- 2 What were the users from Australia?
- 3 What were the users with generated most traffic?

Note

- If the above were questions were known beforehand they would be standing queries.
- Given an application we can optimise it to enable the processing of some kinds of ad-hoc queries.
- In general, it is impossible to be able to accurately answer any possible ad-hoc query.

Issues in Stream Processing

Issues

- Velocity: We may need to give up on processing all data.
- Volume: We may need to build summaries.
 - Not all ad-hoc questions can be answerable.

Possible Solution

- Obtain an approximate answer to the question rather than an exact answer.
- ⇒ Techniques related to [hashing](#) can be very useful.

Programme

- 1 Data Streams
- 2 Sampling and Filtering
 - Sampling Data in a Stream
 - Filtering Streams
- 3 Counting Elements in a Stream

Programme

- 1 Data Streams
- 2 Sampling and Filtering
 - Sampling Data in a Stream
 - Filtering Streams
- 3 Counting Elements in a Stream

Example: Stream of user queries

Suppose we store $1/10$ th of all user queries in order to save space.

- If the ad-hoc query is “how many queries did user u ask?”, the (approximate) answer is easy: $10 \times$ the recorded queries of user u .
- If the ad-hoc query is “how many of u ’s queries were repeated?” the answer is more complicated. Imagine:
 - s is the number of queries recorded once.
 - d is the number of queries recorded twice.
 - No queries were issued more than twice (to simplify this problem).

Then, the (approximate) answer $10 \times d$ is wrong; why?

Working Through the Example

- If the user has made n queries once, $n/10$ will be recorded.
- If the user has made m queries twice, only $m/100$ will be recorded twice (on average).
- Thus, the correct answer is $100 \times d$.

Note

If the ad-hoc query is “what fraction of the user’s queries were repeated?” the answer is more complex; see textbook section 4.2.1, pages 134–135 for a detailed explanation.

Using a Hash Function to Obtain a Representative Sample

- We need to be able to obtain a representative sample of fast stream data.
- Suppose we want to keep the information of $1/10$ th of the users.
- Sometimes, even checking whether a user of a search string is in the list of previous users is too time-consuming.
- By using a hash function we can avoid keeping a list of past users.

Keeping 1/10 of the users — brute force approach

- ① Keep a set of past users u initialised to empty \emptyset .
- ② Keep a set of past selected users s initialised to \emptyset .
- ③ If the user from the incoming stream item is not in u :
 - ① Add the user to u .
 - ② Generate a random number between 0 and 10.
 - ③ If the random number is 0, add the user to s and store the query.
- ④ If the user is in u :
 - ① If the user is also in s , store the query.

But

- The lists u and s could become too large and difficult to maintain!
- Checking whether a user is in u or in s could become too time-consuming!

Keeping 1/10th of the users — using a hash

- 1 Hash user using hash function that maps users to 10 buckets.
- 2 If resulting bucket is 0, store the query.

Notes

- We use hash functions to approximate sampling.
- The resulting approach is much simpler.
- The resulting approach is much faster!

The General Sampling Problem

- The stream consists of tuples with n components.
- A subset of n are the **key** components.
- We want to make a selection of the key components.
- We can use hash functions to obtain a sample consisting of any rational function a/b .
 - 1 Hash the key into b buckets.
 - 2 If the hash value is less than a , record the tuple.
- If the key has more than 1 component, the hash needs to obtain a value for the combined key components.

The General Sampling Problem: Example

The Problem

- A stream of social media posts issues pairs of two components:
 - 1 User ID.
 - 2 Text of the social media post.
- How do we keep samples from two thirds of the users?

The Solution

- 1 Key: User ID
- 2 $a: 2$
- 3 $b: 3$

We hash the user ID into $b = 3$ buckets (0, 1, 2). If the hash value is less than $a = 2$, we record the sample.

The General Sampling Problem: Example

The Problem

- A stream of social media posts issues pairs of two components:
 - 1 User ID.
 - 2 Text of the social media post.
- How do we keep samples from two thirds of the users?

The Solution

- 1 Key: User ID
- 2 a : 2
- 3 b : 3

We hash the user ID into $b = 3$ buckets (0, 1, 2). If the hash value is less than $a = 2$, we record the sample.

Programme

- 1 Data Streams
- 2 Sampling and Filtering
 - Sampling Data in a Stream
 - Filtering Streams
- 3 Counting Elements in a Stream

Example

- We want to allow incoming email only from a whitelist S of authorised users.
- But the list S has 1 billion email addresses!
- We can use hash functions to solve this (again).

Using a hash function

- 1 Define a hash function that maps email addresses into buckets (as many buckets as practical).
- 2 Keep an array with as many bits as hash buckets.
- 3 For every email address in whitelist S , store 1 in the array indexed by the email hash bucket.
- 4 Then, any incoming email address that hashes to a bucket with value 1 stored in the array, is deemed as belonging to whitelist S .

Filtering a Stream, Algorithm

Building the filter

```
hash_filter = [False for i in range(nbuckets)]  
for s in S:  
    hash_filter[hash(s, nbuckets)] = True
```

Testing the filter

```
def in_filter(item):  
    return hash_filter[hash(item, nbuckets)]
```

Questions about Our Example

Questions

- 1 Would this approach generate false positives? (email not in the whitelist is accepted)
- 2 Would this approach generate false negatives? (email in the whitelist is filtered out)
- 3 Which problem is worse? 1 or 2?

The Bloom Filter

- A Bloom filter consists of:
 - 1 An array of n bits, initially all 0's.
 - 2 A collection of k hash functions h_1, h_2, \dots, h_k .
 - Each hash function maps key values to n buckets.
 - 3 A set S of m key values.
- To initialise the bit array:
 - 1 Begin with all bits 0.
 - 2 For each key value v in S and for each hash function h_i :
 - 1 Set 1 to array bit indexed by $h_i(v)$.
- To test a key value w that arrives in the stream:
 - 1 If all $h_1(w), h_2(w), \dots, h_k(w)$ are 1's in the bit array, let the stream element through.
 - 2 If one or more of these bits are 0, reject the stream element.

Bloom Filter, Algorithm

Building the filter

```
hash_filter = [False for i in range(nbuckets)]  
for s in S:  
    for k in range(K):  
        hash_filter[hash((s, k), nbuckets)] = True
```

Testing the filter

```
def in_filter(item):  
    result = True  
    for k in range(K):  
        result = result and hash_filter[hash((item, k),  
                                              nbuckets)]  
    return result
```

Analysis of Bloom Filtering

- If a key value is in S , then the element will pass through the Bloom filter.
- If the key value is not in S , it might still pass (a **false positive**).

What is the probability of a false positive?

(see text book, section 4.3.3, pages 138-139, for an explanation)

- The probability of a false positive is $(1 - e^{(-km/n)})^k$.

Programme

- 1 Data Streams
- 2 Sampling and Filtering
- 3 Counting Elements in a Stream
 - Counting Distinct Elements
 - Exponentially Decaying Windows

Programme

- 1 Data Streams
- 2 Sampling and Filtering
- 3 Counting Elements in a Stream
 - Counting Distinct Elements
 - Exponentially Decaying Windows

The Count-Distinct Problem

The Problem

Suppose you want to count the number of distinct items in a stream, either from the beginning or from some known time in the past.

Issues

- 1 We cannot store all distinct items so far.
- 2 Even using a hash table, the size of the table is limited.

Question

How would you count distinct items using a hash table?

The Flajolet-Martin Algorithm

- The Flajolet-Martin algorithm **estimates** the count of distinct items without keeping track of all past distinct items.
- The count is estimated in an unbiased way.
- The algorithm limits the probability that the estimation error is large.
- The algorithm uses hash functions but it does not need to keep hash tables.

The Flajolet-Martin Algorithm

- 1 Use hash functions that map each of the N elements to at least $\log_2 N$ bits.
 - For example, 2^{64} buckets (64 bits) are enough to hash URL's.
- 2 Define the **tail length** of hash h and item a as the number of trailing zeroes in $h(a)$.
 - If $h(x) = 11010$, then the tail length is 1.
 - If $h(x) = 01000$, then the tail length is 3.
- 3 Let R be the largest tail length observed so far.
- 4 The estimated number of distinct elements is 2^R .

Why it Works: Intuition

- $h(a)$ hashes with equal probability to any of N values.
Therefore:
- About 50% (1 out of 2) of a 's hash to $***0$.
- About 25% (1 out of 2^2) of a 's hash to $**00$.
- About 12.5% (1 out of 2^3) of a 's hash to $*000$.
- So, if we saw the longest tail end = i , then we have probably seen about 2^i distinct items so far.

Combining Several Hash Functions

- Perhaps, by bad luck, a seen object hashes to a bucket with a very high tail end.
- To limit the probability of error, we can combine the tail ends of multiple independent hashes.
 - The **average** is not reliable since it is sensitive to very large numbers.
 - That's why, for example, real estate agents use the median to compare house prices among districts.
 - The **median** would only produce numbers that are powers of 2.
 - The solution is to take the **median of the averages**.
 - Partition the K hashes into groups of G hashes each.
 - Compute the average of the values in each group.
 - Compute the median of the averages.

Programme

- 1 Data Streams
- 2 Sampling and Filtering
- 3 Counting Elements in a Stream
 - Counting Distinct Elements
 - Exponentially Decaying Windows

The Problem of Most-Common Elements

- Suppose we have a stream whose elements are the movie name and the number of tickets sold in the current week.
- We want to find out what are the most popular movies “currently”.
- Which of these movies is more popular “currently”?
 - Movie 1 sold n tickets each week for the last 5 weeks.
 - Movie 2 sold $2n$ tickets last week but none in the previous weeks.
 - Movie 2 sold $10n$ tickets last year but none in the last weeks.
- Depending on what we mean by “currently” we may have a different answer.
- With exponentially decaying windows, we give more importance to the most recent items.

Definition of a Decaying Window

- A decaying window keeps a smooth aggregation of all the counts from the beginning of the stream.
- More recent counts are given more importance.
- If the stream is a_1, a_2, \dots , then the sum of all values with an exponentially decaying window at time t is:

$$\sum_{i=1}^t a_i (1 - c)^{t-i}$$

where c is a fixed constant, e.g. 10^{-6} or 10^{-5} .

Update when a New Item Arrives

- If we already know $s = a_t + a_{t-1}(1 - c) + a_{t-2}(1 - c)^2 + \dots$
- Then we can easily update s when a new a_{t+1} arrives.
- We simply compute $s \leftarrow s(1 - c) + a_{t+1}$.

Proof

$$\begin{aligned} a_{t+1} + s(1 - c) &= \\ a_{t+1} + (a_t + a_{t-1}(1 - c) + a_{t-2}(1 - c)^2 + \dots)(1 - c) &= \\ a_{t+1} + a_t(1 - c) + a_{t-1}(1 - c)^2 + a_{t-2}(1 - c)^3 + \dots \end{aligned}$$

Example: Finding the Most Popular Elements

- Suppose we want to keep counts of the tickets of the most popular movies currently.
- We will generate a stream per movie with 1 each time a ticket for that movie appears in the stream, and 0 otherwise.
- We set c and a threshold of "importance", say 0.5.
- We will keep counts of tickets for those movies whose score is greater than 0.5.
- When a new ticket arrives to the stream:
 - 1 For each movie whose score we are maintaining, multiply its score by $(1 - c)$.
 - 2 If the new ticket is for a movie M :
 - 1 If are maintaining the score for M , add 1 to that score.
 - 2 If there is no score for M , create one and add 1 to that score.
 - 3 If any score falls below the threshold 0.5, drop that score.

How many movies are we maintaining?

$$s = \sum_{i=1}^t a_i (1 - c)^{t-i}$$

- The sum over all weights in an infinite stream is $\sum_{t=0}^{\infty} (1 - c)^t = 1/c$.
- Thus, there cannot be more than $2/c$ movies with score $1/2$ or more.
- In practice, the number of movies with score $1/2$ or more is much less than $2/c$.
- So we can adjust c and the threshold to determine the maximum number of movies we want to keep track at any time.

Sliding versus Decaying Windows

Sliding Window

- Keeps summaries of the last N elements.
- All elements have the same weight.
- We need to worry about the element that falls out of the window when we update the summaries.

Decaying Window

- Keeps summaries of the last N elements.
- More recent elements have higher weight.
- Easy to update the summaries.

Take-home Messages

- The Stream Data Model.
- Sampling of streams.
- Bloom filters.
- Counting distinct elements.
- Counting in the last N elements.

What's Next

Week 9

- Catch-up tutorial on Monday (see announcement in iLearn)
- Assignment 3 ready
- Link Analysis